

PART A
Implement logistic regression using gradient descent in Python.
Once you have written the code, test your model by fitting it to the dataset g
iven below.
For this stage, do not implement regularization, cross validation, the cost fu
nction or use the sklearn library until later.
The purpose of this part of the coursework is to get your logistic regression
code running and to validate your model's results with a small dataset with a k
nown solution.
Refer to the lecture slides on Logistic Regression to see the solution you sho
uld be obtaining with your model.

```
import matplotlib.pyplot as plt
import numpy as np
import math

X = np.array([[0.50],[0.75],[1.00],[1.25],[1.50],[1.75],[1.75],[2.00],[2.25],[2.
50],[2.75],[3.00],[3.25],[3.50],[4.00],[4.25],[4.50],[4.75],[5.00],[5.50]])
Y = np.array([0,0,0,0,0,0,1,0,1,0,1,0,1,0,1,1,1,1,1,1])

def logisticFunc(theta0, theta1, x):
    return float(1) / (1 + np.exp(-(theta0 + theta1 * x)))

def gradientDescent(a):
    tempTheta0 = 0
    tempTheta1 = 0
    theta0 = 0
    theta1 = 0
    iterations = 0

    while True:
        T0_sum = 0
        T1_sum = 0
        for i1 in range(0, X.size):
            T0_sum = (T0_sum + ((logisticFunc(tempTheta0, tempTheta1, X[i1]) - Y[
i1]) * 1))
            T1_sum = (T1_sum + ((logisticFunc(tempTheta0, tempTheta1, X[i1]) - Y[
i1]) * X[i1]))

        tempTheta0 = (tempTheta0 - (a * T0_sum))
        tempTheta1 = (tempTheta1 - (a * T1_sum))

        diffT0 = abs(tempTheta0-theta0)
        diffT1 = abs(tempTheta1-theta1)

        if(iterations > 5000):
            print "Theta0 & Theta1 = ", theta0, theta1
            print "Number of iterations till convergence = ", iterations
            print "Value of alpha = ", a
            return theta0, theta1

        if((diffT0 < 0.0001) and (diffT1 < 0.0001)):
            print "Theta0 & Theta1 = ", theta0, theta1
            print "Number of iterations till convergence = ", iterations
            print "Value of alpha = ", a
            return theta0, theta1

        iterations += 1
        theta0 = tempTheta0
        theta1 = tempTheta1
```

```
T0, T1 = gradientDescent(0.09009)
```

```
Theta0 & Theta1 = [-4.07773037] [ 1.50459548]
```

```
Number of iterations till convergence = 2712
```

```
Value of alpha = 0.09009
```

```

# PART B
# Compute and save the cost function for each iteration of your gradient descent
  algorithm.
# Plot the cost function over the iteration number to answer the following quest
  ions.

import matplotlib.pyplot as plt
import numpy as np

X = np.array([[0.50],[0.75],[1.00],[1.25],[1.50],[1.75],[1.75],[2.00],[2.25],[2.
50],[2.75], [3.00],[3.25],[3.50],[4.00],[4.25],[4.50],[4.75],[5.00],[5.50]])
Y = np.array([0,0,0,0,0,0,1,0,1,0,1,0,1,0,1,1,1,1,1,1])

# Logistic
def logisticFunc(theta0, theta1, x):
    return float(1) / (1 + np.exp(-(theta0 + theta1 * x)))

# Cost function
def costFunc(theta0, theta1, x, y):
    h = logisticFunc(theta0, theta1, x)
    return -y*np.log(h) - (1-y)*np.log(1-h)

def gradientDescent(a):
    tempTheta0 = 0
    tempTheta1 = 0
    theta0 = 0
    theta1 = 0
    iterations = 0
    costIter = np.zeros(0)

    while True:
        costSum = 0
        T0_sum = 0
        T1_sum = 0

        for i1 in range(0, X.size):
            costSum += costFunc(tempTheta0, tempTheta1, X[i1], Y[i1])
            T0_sum = (T0_sum + ((logisticFunc(tempTheta0, tempTheta1, X[i1]) - Y[
]) * 1))
            T1_sum = (T1_sum + ((logisticFunc(tempTheta0, tempTheta1, X[i1]) - Y[
]) * X[i1]))

            tempTheta0 = (tempTheta0 - (a * T0_sum))
            tempTheta1 = (tempTheta1 - (a * T1_sum))

            diffT0 = abs(tempTheta0-theta0)
            diffT1 = abs(tempTheta1-theta1)

        if(iterations > 5000):
            print "Theta0 & Theta1 = ", theta0, theta1
            print "Does not converge before iteration = ", iterations
            print "Value of alpha = ", a
            return theta0, theta1, costIter, iterations

        if((diffT0 < 0.0001) and (diffT1 < 0.0001)):
            print "Theta0 & Theta1 = ", theta0, theta1
            print "Number of iterations till convergence = ", iterations
            print "Value of alpha = ", a
            return theta0, theta1, costIter, iterations

```

```

theta0 = tempTheta0
theta1 = tempTheta1
costIter = np.append(costIter, costSum)
iterations += 1

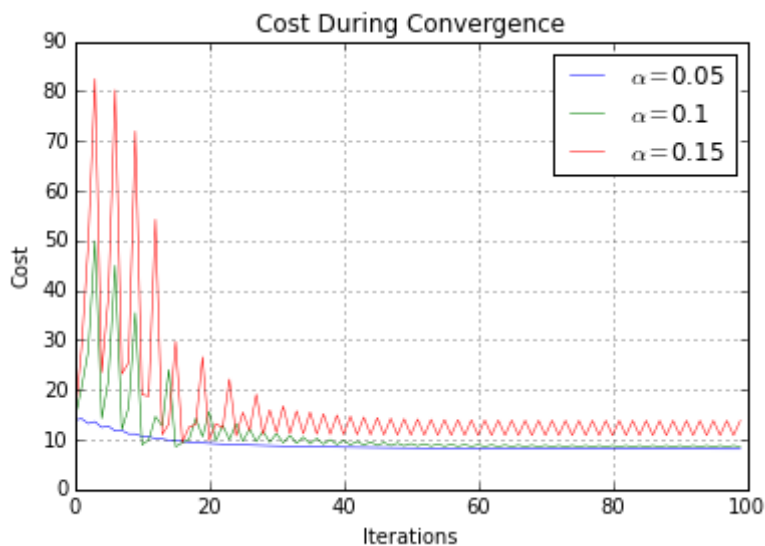
plt.title('Cost During Convergence')
plt.xlabel('Iterations')
plt.ylabel('Cost')

plt.grid(True)

for i in range (0, 3):
    alpha = (i+1) * 0.05
    T0, T1, costIter, iterations = gradientDescent(alpha)
    plt.plot(np.arange( 100 ), costIter[:100], linewidth = 0.5, label=r'$\alpha$'
    =${s}'%(alpha))

plt.legend(loc='upper right')
plt.show()

```



In [1]:

```
# PART C
# Use the sklearn library function to validate your results.
# The function linear_model.SGDClassifier implements gradient descent and the function argument loss = 'log' applies logistic regression.

import matplotlib.pyplot as plt
import numpy as np
from sklearn import linear_model

X = np.array([[0.50],[0.75],[1.00],[1.25],[1.50],[1.75],[1.75],[2.00],[2.25],[2.50],[2.75],[3.00],[3.25],[3.50],[4.00],[4.25],[4.50],[4.75],[5.00],[5.50]])
Y = np.array([0,0,0,0,0,0,1,0,1,0,1,0,1,0,1,1,1,1,1,1,1])

theta0 = -407
theta1 = 150
numIter = 1

while(True):

    clf = linear_model.SGDClassifier(loss='log', n_iter=numIter)
    clf.fit(X, Y)

    if(numIter > 10000):
        print "Model did not converge within 10000 iterations"
        break

    if(round(clf.intercept_*100) == theta0 and round(clf.coef_*100) == theta1):
        break

    numIter+= 1

print "Number of Iterations Required for Convergence: ", numIter

print "Theta0: ", clf.intercept_
print "Theta1: ", clf.coef_

print "Model prediction for 5 hours of study: ", clf.predict(5)
```

```
Number of Iterations Required for Convergence: 4782
Theta0: [-4.07341959]
Theta1: [[ 1.49918078]]
Model prediction for 5 hours of study: [1]
```

In []: